

# Class 8: Breast Cancer Mini Project

James Garza (PID: A16300772)

Before we get stuck into project work we will have a quick look at applying PCA to some example RNASeq data (tail end of lab 7)

Read the data (detailed in lab 7):

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638


Q. How many genes are in this dataset?

```
nrow(rna.data)
```

```
[1] 100
```

```
## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)

## Simple unpolished plot of pc1 and pc2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```



```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642
Cumulative Proportion	0.9262	0.9493	0.96045	0.97152	0.97928	0.98609	0.99251

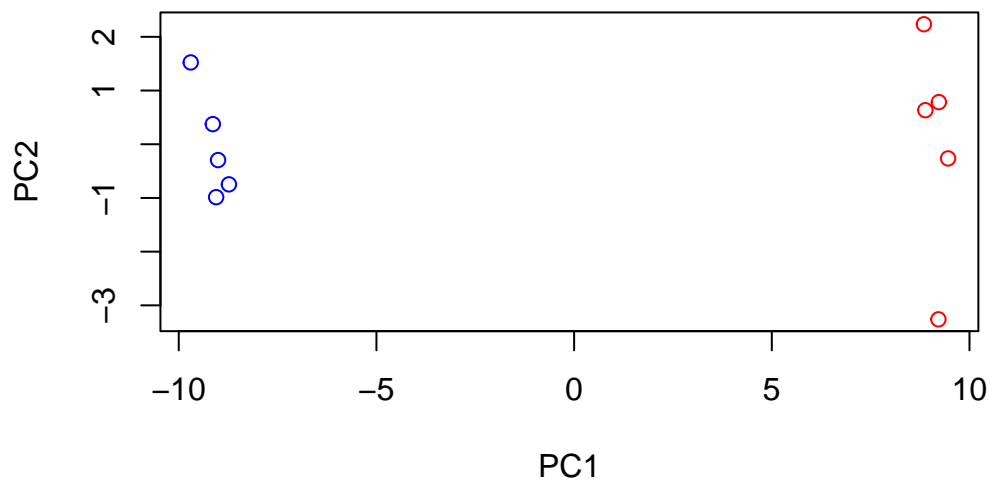
  

	PC8	PC9	PC10
Standard deviation	0.62065	0.60342	3.457e-15
Proportion of Variance	0.00385	0.00364	0.000e+00
Cumulative Proportion	0.99636	1.00000	1.000e+00

```
# WE have 5 wt and 5 ko samples
mycols <- c(rep("blue",5), rep("red", 5))
mycols
```

```
[1] "blue" "blue" "blue" "blue" "blue" "red" "red" "red" "red" "red"
```

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", col=mycols)
```



I could examine which genes contribute most to this first PC

```
pca$rotation[,1]
```

gene1	gene2	gene3	gene4	gene5	gene6
-0.103666005	0.103514749	-0.103761385	0.075320862	0.087428334	0.099670829
gene7	gene8	gene9	gene10	gene11	gene12
0.103609009	0.100759370	-0.085460936	0.103783379	-0.103719665	-0.102001924
gene13	gene14	gene15	gene16	gene17	gene18
-0.103399591	0.102478762	0.099993193	0.103598474	-0.103013773	-0.103774699
gene19	gene20	gene21	gene22	gene23	gene24
-0.103390599	0.103121803	0.103787935	0.102725125	0.103681565	-0.098284250
gene25	gene26	gene27	gene28	gene29	gene30
-0.103302326	-0.085745836	-0.103374849	0.103638752	-0.102739689	0.103044435
gene31	gene32	gene33	gene34	gene35	gene36
-0.101768804	-0.100677376	0.103592988	0.103525731	0.102382706	0.103412422
gene37	gene38	gene39	gene40	gene41	gene42
0.100467583	0.102080752	-0.103744482	-0.102003831	0.103716818	-0.098746675
gene43	gene44	gene45	gene46	gene47	gene48
-0.092001819	-0.103504699	0.103840183	-0.096571619	0.103502386	0.103682769
gene49	gene50	gene51	gene52	gene53	gene54
-0.103188532	-0.103743341	-0.103265591	0.102519795	-0.103245619	-0.103584153

	gene55	gene56	gene57	gene58	gene59	gene60
	0.103695870	-0.103783479	-0.103703675	-0.103503980	-0.103607438	0.103805515
	gene61	gene62	gene63	gene64	gene65	gene66
	-0.103308945	0.103713893	0.101468649	-0.078643996	-0.094219475	0.103845454
	gene67	gene68	gene69	gene70	gene71	gene72
	0.103453646	0.103839510	0.048197107	-0.101365212	0.095664760	0.102347342
	gene73	gene74	gene75	gene76	gene77	gene78
	0.102001050	0.101747637	-0.103592371	0.103514464	-0.097473626	-0.100499426
	gene79	gene80	gene81	gene82	gene83	gene84
	-0.103639415	-0.093476477	-0.100659777	-0.103481127	0.066065263	-0.009263882
	gene85	gene86	gene87	gene88	gene89	gene90
	0.103698370	-0.102122719	-0.103448562	0.098226585	0.100038548	0.103777744
	gene91	gene92	gene93	gene94	gene95	gene96
	-0.103698408	0.092534408	0.102950950	-0.102692869	0.102142492	-0.096658194
	gene97	gene98	gene99	gene100		
	0.100787961	0.103837190	0.103649598	0.103870820		

```
head(sort(abs(pca$rotation[,1]), decreasing = T))
```

	gene100	gene66	gene45	gene68	gene98	gene60
	0.1038708	0.1038455	0.1038402	0.1038395	0.1038372	0.1038055

## Analkysis of Breast Cancer FNA (fine needle aspirations) data.

```
# read.csv("https://bioboot.github.io/bimm143_S20/class-material/WisconsinCancer.csv")
```

```
# Save your input data file into your Project directory
fna.data <- "WisconsinCancer.csv"
```

```
# Complete the following code to input the data and store as wisc.df
wisc.df <- read.csv(fna.data, row.names=1)
```

```
# Create diagnosis vector for later
diagnosis <- as.factor(wisc.df$diagnosis)
diagnosis
```

```
[1] M M M M M M M M M M M M M M M M M M B B B M M M M M M M M M M M M M
[38] B M M M M M M M M B M B B B B B M M B M M B B B B M B M M B B B B M B M M
```

```

[75] B M B M M B B B M M B M M M B B B M B B M M B B B M B B M B B
[112] B B B B B B M M M B M M B B B M M B M B M M B M M B B M B B M B B B M B
[149] B B B B B B B B M B B B B M M B M B B M M B B M M B B B M B B M M M B M
[186] B M B B B M B B M M B M M M M B M M M B M B M B B M B M M M M B B M M B B
[223] B M B B B B B M M B B M B B M M B M B B B B M B B B B M B M M M M M M M
[260] M M M M M M M B B B B B B M B M B B M B B M B M M B B B B B B B B B B
[297] B M B B M B M B B B B B B B B B B B B B B M B B B M B M B B B B M M M B B
[334] B B M B M B M B B B M B B B B B B B M M M B B B B B B B B B B M M B M M
[371] M B M M B B B B B M B B B B B M B B B M B B M M B B B B B B M B B B B B B
[408] B M B B B B B M B B M B B B B B B B B B B B M B M M B M B B B B B M B B
[445] M B M B B M B M B B B B B B B B M M B B B B B B M B B B B B B B B B M B
[482] B B B B B B M B M B B M B B B B B M M B M B M B B B B B M B B M B M B M M
[519] B B B M B B B B B B B B B B B M B M M B B B B B B B B B B B B B B B B
[556] B B B B B B M M M M M M B
Levels: B M

```

Now I want to make sure I remove that column from my dataset for analysis

```

# We can use -1 here to remove the first column
wisc.data <- wisc.df[,-1]
head(wisc.data)

```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
842302	17.99	10.38	122.80	1001.0	0.11840
842517	20.57	17.77	132.90	1326.0	0.08474
84300903	19.69	21.25	130.00	1203.0	0.10960
84348301	11.42	20.38	77.58	386.1	0.14250
84358402	20.29	14.34	135.10	1297.0	0.10030
843786	12.45	15.70	82.57	477.1	0.12780

	compactness_mean	concavity_mean	concave.points_mean	symmetry_mean
842302	0.27760	0.3001	0.14710	0.2419
842517	0.07864	0.0869	0.07017	0.1812
84300903	0.15990	0.1974	0.12790	0.2069
84348301	0.28390	0.2414	0.10520	0.2597
84358402	0.13280	0.1980	0.10430	0.1809
843786	0.17000	0.1578	0.08089	0.2087

	fractal_dimension_mean	radius_se	texture_se	perimeter_se	area_se
842302	0.07871	1.0950	0.9053	8.589	153.40
842517	0.05667	0.5435	0.7339	3.398	74.08
84300903	0.05999	0.7456	0.7869	4.585	94.03
84348301	0.09744	0.4956	1.1560	3.445	27.23
84358402	0.05883	0.7572	0.7813	5.438	94.44

843786	0.07613	0.3345	0.8902	2.217	27.19
	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	0.006399	0.04904	0.05373	0.01587	
842517	0.005225	0.01308	0.01860	0.01340	
84300903	0.006150	0.04006	0.03832	0.02058	
84348301	0.009110	0.07458	0.05661	0.01867	
84358402	0.011490	0.02461	0.05688	0.01885	
843786	0.007510	0.03345	0.03672	0.01137	
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	
84300903	0.02250	0.004571	23.57	25.53	
84348301	0.05963	0.009208	14.91	26.50	
84358402	0.01756	0.005115	22.54	16.67	
843786	0.02165	0.005082	15.47	23.75	
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
84358402	152.20	1575.0	0.1374	0.2050	
843786	103.40	741.6	0.1791	0.5249	
	concavity_worst	concave.points_worst	symmetry_worst		
842302	0.7119	0.2654	0.4601		
842517	0.2416	0.1860	0.2750		
84300903	0.4504	0.2430	0.3613		
84348301	0.6869	0.2575	0.6638		
84358402	0.4000	0.1625	0.2364		
843786	0.5355	0.1741	0.3985		
	fractal_dimension_worst				
842302	0.11890				
842517	0.08902				
84300903	0.08758				
84348301	0.17300				
84358402	0.07678				
843786	0.12440				

Q1. How many observations are in this dataset?

```
nrow(wisc.data)
```

```
[1] 569
```

There are 569 patients in this dataset

```
length(wisc.data)
```

```
[1] 30
```

There are 30 columns they are looking at.

Q2. How many of the observations have a malignant diagnosis?

```
table(wisc.df$diagnosis)
```

```
  B    M  
357 212
```

There are 212 malignant individuals

Q3. How many variables/features in the data are suffixed with `_mean`?

```
colnames(wisc.data)
```

```
[1] "radius_mean"      "texture_mean"  
[3] "perimeter_mean"   "area_mean"  
[5] "smoothness_mean"  "compactness_mean"  
[7] "concavity_mean"    "concave.points_mean"  
[9] "symmetry_mean"     "fractal_dimension_mean"  
[11] "radius_se"         "texture_se"  
[13] "perimeter_se"      "area_se"  
[15] "smoothness_se"     "compactness_se"  
[17] "concavity_se"      "concave.points_se"  
[19] "symmetry_se"       "fractal_dimension_se"  
[21] "radius_worst"      "texture_worst"  
[23] "perimeter_worst"   "area_worst"  
[25] "smoothness_worst"  "compactness_worst"  
[27] "concavity_worst"   "concave.points_worst"  
[29] "symmetry_worst"    "fractal_dimension_worst"
```

```
length(grep("_mean", colnames(wisc.data)))
```

[1] 10

There are 10 columns with the name “\_\_mean” in them.

## Principa Component Analysis (PCA)

Here we will use `prcomp()` on the `wisc.data` object - the one without the diagnosis column.

First we have to decide whether to use the `scale=TRUE` argument when we run `prcomp()`. We can look at the means and sd of each column. If they are similar then we are all good to go. IOf not we should use `scale=TRUE`.

```
# Check column means and standard deviations
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data,2,sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01



area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

These are very different so we should `scale=TRUE`.

```
# Perform PCA on wisc.data by completing the following code
wisc.pr <- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997

	PC29	PC30
Standard deviation	0.02736	0.01153
Proportion of Variance	0.00002	0.00000
Cumulative Proportion	1.00000	1.00000

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44.27% is captured by the first PC

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

3 PCs are needed to get 70% of the variance


Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 PCs are needed to get 90% of the variance

Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

What stands out to me is that it is very difficult to understand because it is designed for small datasets and ours is too big.

```
biplot(wisc.pr)
```




We need to make our own plot.

```
attributes(wisc.pr)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"

$class
[1] "prcomp"
```

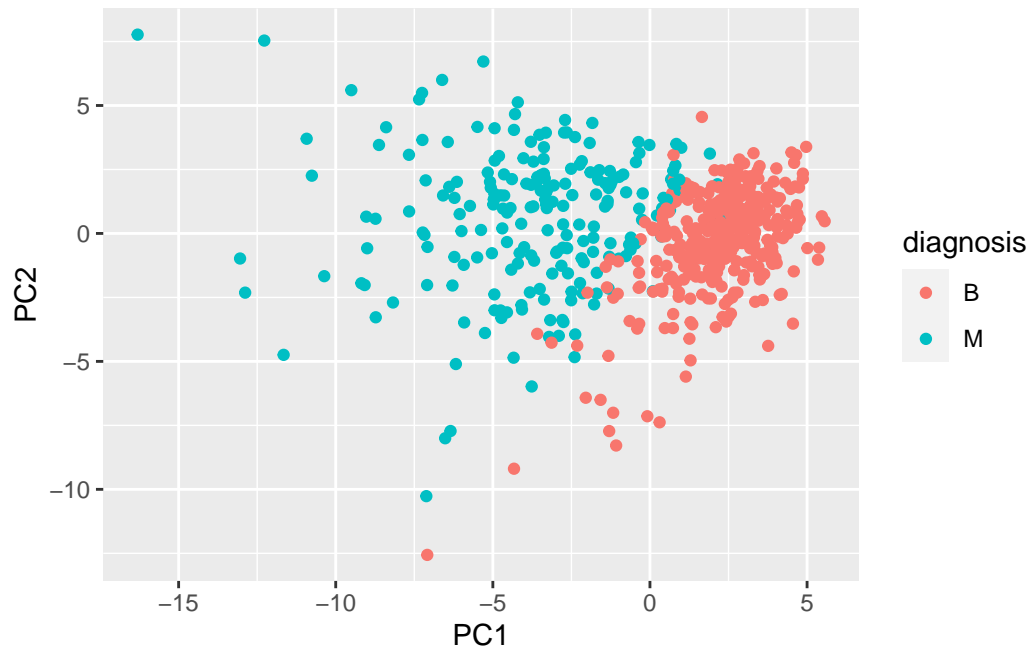
```
plot(wisc.pr$x[,1], wisc.pr$x[,2], col=diagnosis)
```



```
library(ggplot2)

pc <- as.data.frame(wisc.pr$x)

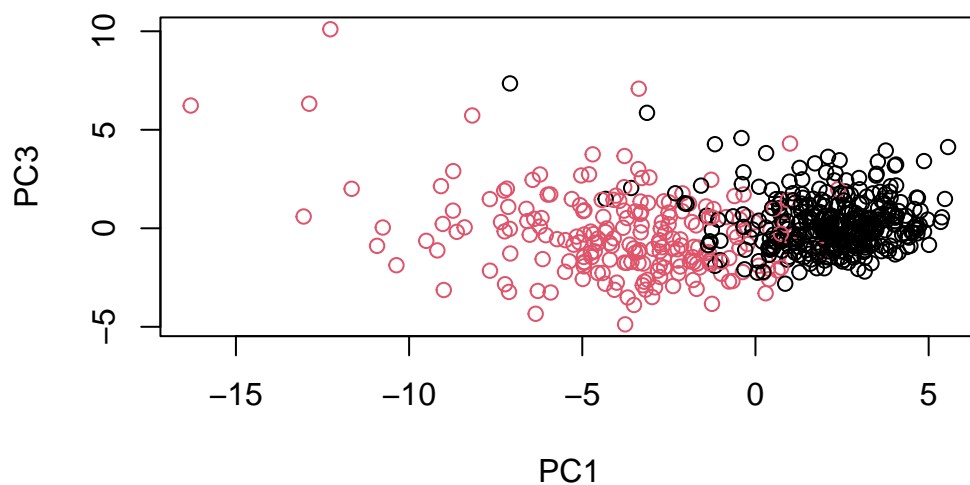
ggplot(pc) +
  aes(PC1, PC2, col=diagnosis)+
  geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

There are still two clear groupings between the benign and malignant samples between PC1 and PC3.

```
# Repeat for components 1 and 3
plot(wisc.pr$x[,1], wisc.pr$x[,3], col = diagnosis,
     xlab = "PC1", ylab = "PC3")
```



Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
wisc.pr$rotation["concave.points_mean",1]
```

```
[1] -0.2608538
```

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

```
tbl <- summary(wisc.pr)
which(tbl$importance[3,] > 0.8)[1]
```

```
PC5
```


```
5
```

5 PCs are needed to explain 80% of the variation in the dataset

## Hierarchical clustering


The main function for Hierarchical clustering is called `hclust()`. It takes a distance matrix as input.

```
d <- dist( scale(wisc.data))
wisc.hclust <- hclust(d)
plot(wisc.hclust)
```



```
plot(wisc.hclust)
abline(h=20, col="purple")
```

## Cluster Dendrogram



d  
hclust (\*, "complete")

```
grps <- cutree(wisc.hclust, h=20)
table(grps)
```

```
grps
  1  2  3  4
177  7 383  2
```

Q11. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

At height 20 there are 4 clusters.

Come back here later to see how our cluster `grps` correspond to M or B groups.

```
table(grps)
```

```
grps
  1  2  3  4
177  7 383  2
```




## Combining Methods

Here we will perform Clustering PCA Results rather than the original data.

In other words we will cluster using `wisc.pr$x` - our new better variables or PCs. We can choose as many or as few PCs to use as we like. It is your call!

```
d.pc <- dist (wisc.pr$x[,1:3] )  
wisc.pr.hclust <- hclust(d.pc, method="ward.D2")  
plot(wisc.pr.hclust)
```


### Cluster Dendrogram



d.pc  
hclust (\*, "ward.D2")

```
plot(wisc.pr.hclust)  
abline(h=80, col="red")
```

## Cluster Dendrogram



d.pc  
hclust (\*, "ward.D2")

```
grps <- cutree(wisc.pr.hclust, h=80)
table(grps)
```

```
grps
  1   2
203 366
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

We can use `table()` function to make a cross-table as well as just a count table.

```
table(diagnosis)
```

```
diagnosis
  B   M
357 212
```

```
table(grps, diagnosis)
```

```

diagnosis
grps  B   M
1    24 179
2   333  33

```

Write a note here about how to read this cross-table result. The results indicate that our cluster 1 mostly captures cancer (M) and our cluster 2 mostly captures our healthy (B) from samples/individuals.

## 7. Prediction

```

#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc

```

```

          PC1      PC2      PC3      PC4      PC5      PC6      PC7
[1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
[2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
          PC8      PC9      PC10     PC11     PC12     PC13     PC14
[1,] -0.2307350  0.1029569 -0.9272861  0.3411457  0.375921  0.1610764  1.187882
[2,] -0.3307423  0.5281896 -0.4855301  0.7173233 -1.185917  0.5893856  0.303029
          PC15     PC16     PC17     PC18     PC19     PC20
[1,]  0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
[2,]  0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
          PC21     PC22     PC23     PC24     PC25     PC26
[1,]  0.1228233  0.09358453  0.08347651  0.1223396  0.02124121  0.078884581
[2,] -0.1224776  0.01732146  0.06316631 -0.2338618 -0.20755948 -0.009833238
          PC27     PC28     PC29     PC30
[1,]  0.220199544 -0.02946023 -0.015620933  0.005269029
[2,] -0.001134152  0.09638361  0.002795349 -0.019015820

```


And plot this up

```

plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)

```

```
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q18. Which of these new patients should we prioritize for follow up based on your results?

We should follow up with patients that are benign but are near cluster 1 because that indicates that they would be best suited for follow ups because they are have similar data to those with malignancy which suggests that they deserve a second look at their respective cancer cases.